

Basics

	Sample statistic	Population Parameter
Mean	\bar{X}	μ
Variance	S^2	σ^2
Correlation	r	ρ
Noise	e_i	ϵ_i
	Guess	True, but unknown

The Mean, Variance and StdDev are subject to outliers.

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\mu = E(W) = E(a + bX) = a + bE(X)$$

$$\sigma^2 = Var(X) = E[X^2] - \mu_x^2$$

$$S^2 = \sum_{all\ i} \frac{(X_i - \bar{X})^2}{n - 1}$$

$$\sigma = StdDev(X) = \sqrt{\sigma^2}$$

$$Var(X) = E[(X - \mu)^2] = E(X^2) - \mu^2$$

$$= \sum_{all\ x} (x - \mu_x)^2 P(X = x)$$

$$= E(X^2) - E(X)^2 = E(X - E(X))^2$$

$$Var(X + c) = Var(X)$$

$$Var(cX) = c^2 Var(X)$$

$$Var(\bar{X}) = \frac{\sigma^2}{n}$$

If X,Y are dependent

$$Var(X + Y) \neq Var(X) + Var(Y)$$

$$Var(X + Y) = Var(X - Y)$$

$$Var((a + bX) + (c + dY)) = b^2 Var(X) + d^2 Var(Y) + 2bd Cov(X, Y)$$

$$\begin{aligned} Var(aX + bY) &= a^2 Var(X) + b^2 Var(Y) + 2ab Cov(X, Y) \\ &= a^2 Var(X) + b^2 Var(Y) + 2ab \sigma_x \sigma_y \rho_{xy} \end{aligned}$$

Quartiles split the data into 4 equal groups by number of values, or 25% percentiles. Q2 is the median.

Study Design

Types of variables: categorical vs. quantitative, dummy, discrete vs continuous. Observational vs Experiments. Only experimentals can show causal effects. Key factors to an experiment: control group, randomized, and replicated. Usually 30 replications is sufficient. Avoid confounding by having cross-over groups.

Scope of Inference: Interval Validity: low when 1) unaccounted confounding factors, 2) ignored missing data 3) noncompliance 4) unverified assumptions 5) suboptimal method of analysis. Applicable when allocation of units to groups is random. **External Validity:** high when can be generalized to population. Applicable when selection of units is random.

Covariance and Correlation

Covariance gives direction.

Correlation gives direction and strength.

Both are *linear*.

$$Cov(X, Y) = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

$$Cor = \frac{Cov}{\sigma_x \cdot \sigma_y}$$

Probabilities

$0 \leq$ All probabilities ≤ 1

Mutually exclusive and Exhaustive

$$P(\bar{A}) = 1 - P(A)$$

$$P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$$

$$P(A|B) = \frac{P(A \text{ and } B)}{P(B)}$$

$$P(B) = \sum_{i=1}^n P(B \cap A_i) = \sum_{i=1}^n P(B|A_i)((A_i)$$

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

Joint vs. Marginal probabilities

Independent if $P(A|B) = P(A)$

For Independent only: $P(E \text{ and } F) = P(E) \cdot P(F)$

	B	\bar{B}
A	$P(A \text{ and } B) = P(B)P(A B)$	$P(A \text{ and } \bar{B}) = P(\bar{B})P(A \bar{B})$
\bar{A}	$P(\bar{A} \text{ and } B) = P(A)P(B \bar{A})$	$P(\bar{A} \text{ and } \bar{B}) = P(\bar{A})P(\bar{B} \bar{A})$

all success	p^n
all failure	$(1 - p)^n$
at least one failure	$1 - p^n$
at least one success	$1 - (1 - p)^n$

Random Variables

$$P(X \leq x) \rightarrow CDF$$

$$E(cX) = c \cdot E(X)$$

$$E(X + c) = E(X) + c$$

$$E(X + Y) = E(X) + E(Y)$$

If $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ then $\bar{X} \sim N(\mu_{\bar{X}} = \mu, \sigma_{\bar{X}}^2 = \sigma^2/n)$.

If pop dist is not normal, S^2 will eventually be Normal but may not be independent.

Bias of an Estimator

In practice n has to relatively much larger like > 100 .

Guesses should be *unbiased* and have *minimum variance*.

MVUE (Minimum Variance, Unbiased Estimates).

$$bias(\hat{\theta}) = E(\hat{\theta}) - \theta$$

Unbiased if $bias = 0$ (expected value equals true, not a particular value of \bar{x}).

For samples, we divide by $n - 1$ instead of n to make an unbiased estimator. The guess would otherwise be too low.

$$E(X) = \sum_{all\ x} xP(X = x)$$

$$E(\bar{X}) = \sum_{i=1}^{\infty} x_i p_i = \mu$$

$$E(S^2) = \sigma^2$$

$$E[X + c] = E[X] + c$$

$$E[X + Y] = E[X] + E[Y]$$

$$E[aX] = aE[X]$$

$$E[aX + bY + c] = aE[X] + bE[Y] + c$$

$$E((a + bX) + (c_d Y)) = a + bE(X) + c + dE(Y)$$

Example: Roulette has a \$1 bet with a \$35 payoff for $\frac{1}{38}$ odds.

$$E[\text{gain from a \$1 bet}] = -\$1 \cdot \frac{37}{38} + \$35 \cdot \frac{1}{38} = -\$0.0526$$

Uniform Distribution

$$E(X) = \frac{(a + b)}{2}$$

$$Var(X) = \frac{(b - a)^2}{12}$$

y axis should be a fraction to make area = 1

Normal Distribution

$$X \sim \mathcal{N}(\mu, \sigma^2)$$

$$Z = \frac{X - \mu}{\sigma} \sim \mathcal{N}(0, 1)$$

$$X = \sigma Z + \mu$$

$\Phi(z)$ =distribution

$$P(a \leq X \leq b) = P[(a - \mu) \leq (X - \mu) \leq (b - \mu)]$$

$$= P\left[\frac{(a - \mu)}{\sigma} \leq \frac{(X - \mu)}{\sigma} \leq \frac{(b - \mu)}{\sigma}\right]$$

$$= P\left[\frac{(a - \mu)}{\sigma} \leq Z \leq \frac{(b - \mu)}{\sigma}\right]$$

Linear combination of normals is normal. Used by CLT.

Binomial Distribution

- n independent trials
- binary result
- same probability of success
- total number of successes

$$X \sim \text{Bin}(n, p)$$

$$\binom{n}{x} = \frac{n!}{x!(n-x)!}$$

$$P(X = x) = \frac{n!}{x!(n-x)!} p^x q^{(n-x)}$$

$$\mu_x = E(X) = n \cdot p$$

$$\sigma^2 = \text{Var}(X) = n \cdot p \cdot q = n \cdot p \cdot (1-p)$$

$$X \sim N(\mu = np, \sigma = \sqrt{np(1-p)})$$

$$\hat{p} \sim N(\mu = p, \sigma \sqrt{p(1-p)/n})$$

Shape of distribution depends on p, n .

Small p , left skewed. Large p , right skewed

Hypothesis Test - General

A **scientific hypothesis** makes a testable statement about the observable universe.

A **statistical hypothesis** is more restricted. It concerns the behavior of a measurable (or observable) random variable. It is often a statement or claim about a parameter of a population or distribution.

1. Formulate hypotheses (H_0, H_a, α)
2. Calculate test statistic and reference distrib.
3. Calculate p -value based on reference distribution of test statistic $|H_0$. **p value is NOT the probability that the null hypothesis is true.** It is probability of seeing our data on the null hypothesis.
4. Determine conclusion and scope of inference

Note: t values are measures of *distance*. p values are measure of *probability*. Assumptions? Independence? Distribution of Observations? Parametric vs. non-parametric?

Hypotheses	Decision Rule
$H_0 : \mu = \mu_0$ $H_a : \mu \neq \mu_0$	If $ t_{stat} > 1.96$, reject H_0
$H_0 : \mu = \mu_0$ $H_a : \mu < \mu_0$	
$H_0 : \mu = \mu_0$ $H_a : \mu > \mu_0$	If $t_{stat} < -1.64$, reject H_0
$H_0 : \mu = \mu_0$ $H_a : \mu > \mu_0$	If $t_{stat} > 1.64$, reject H_0

Assumptions:

1. Observations are **independent**. This is true if the sample is randomly selected, but false if there is bias in the sampling.

2. **Normal Distributions** of the observations. This happens when the sample is large enough (via the Central Limit Theorem) or it is known that the population is normally distributed.

If these assumptions are not correct, none of the t -tests will work correctly.

Hypothesis Test - Mean

Calculation by hand using a t test:

$$t_{stat} = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$$

Fisher's Randomization Test

For casual inference in experiments. No assumption as to the underlying distribution has to be made. Generally, is $\mu_1 = \mu_2$? The only assumption is Additive Treatment Effect (δ):

$$Y_{c,i} = Y_{v,i} + \delta$$

$H_0 : \delta = 0$, ie zero treatment effect for all units. $H_a : \delta \neq 0$ non-zero treatment effect for ALL units. Reference distrib built through simulation. Non-parametric.

Assumptions:

- Random assignment to groups
- Under H_0 , *independence*, *interchangability* of study units.

Test Statistic: Difference of outcomes of the two groups: (could also use difference between the medians)

$$\hat{\delta} = \bar{Y}_c - \bar{Y}_v$$

The maximum number of possible simulations is the binomial distribution:

$$\binom{t}{s}$$

where t is the total number of study units and s is the number where the effect was shown.

Calculating p value: The p value is proportion of values above or below the observed test statistic.

$$\hat{\delta} = 0.5$$

$$p = P(\bar{Y}_c - \bar{Y}_v \geq \bar{y}_c - \bar{y}_v)$$

$$= P(\bar{Y}_c - \bar{Y}_v \geq 0.5) = 0.029$$

so there is sufficient evidence to reject the null.

Scope of inference: Internal Validity maybe not if the same hospital. External validity: possibly not, as these were volunteers.

Permutation Test

For observational studies, comparing two groups - for generalization. Means of observations are presumed the same $E(Y_{1,i}) = \mu_1, E(Y_{2,j}) = \mu_2, H_0 : \mu_1 = \mu_2$.

Test stat: $\hat{D} = \bar{Y}_i - \bar{Y}_1$. Reference distrib by simulation. Assumes independence of study units and groups, nonparametric, similar shapes and spreads.

All Things χ^2

$$\chi_1^2 = Z^2$$

$$\chi_n^2 = \sum_{i=1}^n Z_i^2$$

$$t_n = \frac{Z}{\sqrt{\chi_n^2/n}}$$

$$F_{n,m} = \frac{\chi_n^2/n}{\chi_m^2/m}$$

$$\chi_{n+m}^2 = \chi_n^2 + \chi_m^2$$

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{df=n-1}^2$$

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{x})^2 \sim \frac{\sigma^2}{n-1} \chi_{df=n-1}^2$$

The t distribution

The t distribution has a slightly wider spread in the tails than Standard Normal. Also known as student-t distribution. This happens when you don't know the real σ and are forced to use S , which is less reliable.

$$\bar{X} \approx N\left(\mu, \frac{\sigma^2}{n}\right)$$

allows us to Z-score

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$$

but we don't know σ so in practice we use

$$T = \frac{\bar{X} - \mu}{S/\sqrt{n}} \sim t_{n-1}$$

(which has two random variables (\bar{X} and S)). Now, take that previous definition of T and add σ/\sqrt{n} to the top and bottom:

$$T = \frac{\bar{X} - \mu}{S/\sqrt{n}}$$

$$= \frac{(\bar{X} - \mu)/(\sigma/\sqrt{n})}{(S/\sqrt{n})/(\sigma/\sqrt{n})}$$

$$= \frac{Z}{\sqrt{S^2/\sigma^2}}$$

$$= \frac{Z}{\sqrt{\frac{(n-1)S^2}{(n-1)\sigma^2}}}$$

$$= \frac{Z}{\sqrt{\chi^2/df}}, \text{ since } \chi^2 = \frac{(n-1)S^2}{\sigma^2}$$

Assumptions: 1) Observations are independent (randomly selected), 2) Normal distrib of observations (either > 30 via CLT or pop. is known $\sim N(\mu, \sigma^2)$).

Confidence Intervals

CI is Estimate \pm margin of error or Estimate $\pm (z \text{ value}) \times$ (SD of estimate).

$$\mu_0 = \bar{X} \pm t \cdot s / \sqrt{n}$$

If the interval spans 0, then one is not significantly bigger than the other (or cannot be determined). As long as $n > 30$, it doesn't matter if the sample size is different between the random variables.

To find the 95% Confidence Interval for the **true mean** μ :

$$\bar{x} \pm t_{df=n-1}^* (s / \sqrt{n})$$

t^* has to be looked up based on the reference distrib.

Types of Errors

Type I the null hypothesis is rejected when it is true

Type II the null hypothesis is accepted when it is false

Power the probability of correctly rejecting the null when the alternative is actually true.

α is *level of significance* - probability of making a Type I error. The greater the cost of an error, the smaller α should be. β is the probability of making a Type II error. There is an inverse relation between Type I and II errors. Reducing one increases the other. The only way to reduce both is to increase n the sample size.

Pooled Two Sample t -Test

assumes equal variances, different means, all observations are independent, random, > 30 or known normal distrib.

$$H_0 : \mu_1 - \mu_2 = 0$$

$$T = \frac{(\bar{x}_1 - \bar{x}_2) - (\mu_1 - \mu_2 | H_0)}{S_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

$$T \sim t_{df=n_1+n_2-2}$$

$$CI : (\bar{x}_1 - \bar{x}_2) \pm t^* \left(S_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}} \right)$$

$$S_p^2 = \frac{(n_1 - 1)S_1^2 + (n_2 - 1)S_2^2}{n_1 + n_2 - 2}$$

Do not use if $\frac{S_1^2}{S_2^2} \geq 2$.

Unpooled Two Sample t -Test

assumes unequal variances, all observations independent, random, > 30 or known normal distrib.

$$H_0 : \mu_1 - \mu_2 = 0$$

$$T = \frac{(\bar{x}_1 - \bar{x}_2) - (\mu_1 - \mu_2 | H_0)}{\sqrt{\frac{S_1^2}{n_1} + \frac{S_2^2}{n_2}}}$$

$$T \sim t_{df=\min(n_1, n_2)-1}$$

$$CI : (\bar{x}_1 - \bar{x}_2) \pm t^* \left(S_p \sqrt{\frac{S_1^2}{n_1} + \frac{S_2^2}{n_2}} \right)$$

Advantage Pooling: slightly more power to detect a difference. Disadvantage pooling: extra assumption could be wrong, with a high type I error rate.

Paired Test

This when there are two samples that are **not** independent, e.g. Weight Watchers, Before / After or matched, shared characteristics. Is the data matched or independent? If we don't take into account the match, the results are wrong. Applies when there are two obs for each study unit. To account for this, take the difference between $\bar{X}_1 - \bar{X}_2$ and then do a hypothesis test on the *difference*.

Basically it is a one-sample test on the differences.

$$H_0 : \mu_D = 0. T \sim t_{df=n_D-1}$$

Comparing Two Proportions

(Not really covered in this class, but may be useful).

$$Var(\hat{p}_1 - \hat{p}_2) = \frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}$$

The 95% confidence interval for $p_1 - p_2$ is:

$$(\hat{p}_1 - \hat{p}_2) \pm 1.96 \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

Decision Rules for Testing Two Proportions:

$$T = \frac{(\hat{p}_1 - \hat{p}_2)}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}, \text{ where } \hat{p} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$$

\hat{p} is called the *pooled proportion*.

Comparing Two Normal Distributions

$$A \sim \mathcal{N}(\mu_A, \sigma_A^2)$$

$$B \sim \mathcal{N}(\mu_B, \sigma_B^2)$$

Find $P(A < B + c)$ then Z-score:

$$P(A - B - c < 0)$$

$$(A - B - c) \sim \mathcal{N}(\mu_A - \mu_B - c, \sigma_A^2 + \sigma_B^2 - 2Cov(A, B))$$

$$P(X = A - B - c > 0)$$

Assumptions

Robustness: able to be useful when assumptions are violated. Check for violations graphically (QQplot, box plot, density) and look at the study design. Under the null and correct assumptions, p -values should be uniformly distributed. t -test is very robust when sample sizes are large, or both groups have the same size. Not so good when one group is very skewed. Fix with unpooled t test, data transformations, or non-parametric tests.

F -Distribution and F -tests

$$R = \frac{X/n_x}{Y/n_y} \sim F_{n_x, n_y}$$

Hypothesis: $H_0 : \sigma_x^2 = \sigma_y^2, H_a : \sigma_x^2 / \sigma_y^2 \neq 1$.

Test statistic:

$$F = \frac{S_X^2}{S_Y^2} \sim F_{n_x-1, n_y-1}$$

Transformations

Purpose: make data more symmetric

If Right skewed, use: $1/Y, \log(Y), \sqrt{Y}$. Left skewed: $Y^2, e^Y, (n - Y)$

If done on log scale: 1) be cautious about ≤ 0 . Use ratios of medians as $\exp(\text{mean})$ will not be useful. Consider rank transform.

Rank Sum Test

non-parametric, 2 indep. groups. Rank ALL data (2 groups combined) and sum up the ranks in one of the groups. Average ties. H_0 : medians are the same. Reference distrib: small n simulate like permutation test; large n approx normality if $n_j \geq 10$:

$$T \sim N(n_x \bar{R}, S_R^2 \frac{n_x n_y}{n_x + n_y})$$

Sign Test - Paired Data

aka binomial test, nonparametric. Take all the differences

Hypotheses: $H_0 : P(D_i > 0) = P(Y_i > X_i) = 0.5$.

Test stat: $K = \sum_{i=1}^n I(D_i > 0)$. ie. count up the number of times the difference is > 0 .

Reference dist: Small n : $K \sim \text{Binom}(n, 0.5)$. Large n : $K \sim N(m/2, m/4)$. Large is considered $np > 10$ and $n(1 - p) > 10$.

Signed Rank Test

Hypotheses: H_0 : rank of magnitude of within pair differences is unrelated to the sign of the difference.

$D_i = Y_i - X_i$, and the calc Rank($|D_i|$) for all n pairs.

Test stat: $W = \sum_{i=1}^n \text{Ranks}_{i|D_i > 0}$.

Reference dist: small n : use permutation. Large n : $W \sim N(m(m+1)/4, m(m+1)(2m+1)/24)$

m is the number of non-zero observations. Like the sign test, **we throw away the zeros**.

Power

Power increases with: larger n , moving alternate distrib further away, smaller σ or S .

$$n \geq \left(\frac{(z^*)(\sigma)}{m} \right)^2$$
$$n \geq \left(\frac{\sigma(|z_{\alpha/2}| + |z_{\beta}|)}{\mu_a - \mu_0} \right)^2$$

Bonferroni: $\alpha^* = \frac{\alpha}{\text{tests}}$. Prob of error: $1 - (1 - \alpha)^I$

ANOVA

Comparing 3 or more groups. F -stat to test H_0

$$F = \frac{MSB}{MSW} = \frac{SSB/df_B}{SSW/df_W} = \frac{SSB/(I-1)}{SSW/(N-I)}$$

Errors between groups: how far away is the variance of one group from all the other groups?

$$SSM = SSB = \sum_{i=1}^I \sum_{j=1}^{n_i} (\bar{Y}_i - \bar{Y})^2 = \sum_{i=1}^I n_i (\bar{Y}_i - \bar{Y})^2$$

with degrees of freedom $df_M = I - 1$.

Errors within groups: how far away is the variance within the group?

$$SSE = SSW = \sum_{i=1}^I \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_i)^2 = \sum_{i=1}^I (n_i - 1) S_i^2$$

with degrees of freedom: $df_E = n - I$

There are two ways of determining which group differ. There are a number of techniques. 1) measure all pair-wise combinations, 2) use contrasts.

Steps in a Complete ANOVA Procedure:

1. Examine the data, checking assumptions
2. If possible, formulate in advance some working (alternative) hypotheses about how the population group means might differ. For example, if we are examining a control group vs two study groups, use

$$H_0 : \mu_1 - \frac{1}{2}(\mu_2 + \mu_3) = 0$$

This is called a **contrast**, which is a linear combination of μ 's. The coefficients sum up to 0. If the coefficients did not sum up to 0, then the null hypothesis would not hold. If the groups had different sample sizes, you could make the coefficient a weighting factor to adjust for this.

3. Check the evidence against the global null hypothesis of no differences among the groups by calculating the F -test, or by testing all pairwise combinations.
4. If the F -test is significant (that is, if it leads to a rejection of the null hypothesis of equal population group means):

- Test the individual hypotheses specified at the second step

- If there was not enough information to formulate working hypotheses, test all pairwise comparisons of means, adjusting for multiple comparisons (example also coming)

Source	SS	DF	MS	F
Groups (between)	SSB	$I - 1$	$S_B^2 = SSB/DFB$	$F = MSB/MSW$
Error (within)	SSW	$n - I$	$S_W^2 = SSW/DFW$	
Total	SST	$n - 1$	SST/DFT	

Contrast Test Hypothesis

This is similar to a pooled test (we are assuming the variance among each group is the same):

$$H_0 : \psi = 0$$

$$H_a : \psi \neq 0$$

$$T = \frac{\sum a_i \bar{Y}_i}{S_p \sqrt{\sum \frac{a_i^2}{n_i}}}$$

This t -test has a t distribution with degrees of freedom associated with S_p under H_0 . The test can be 1- or 2-sided. For this t -test, what is the variance? If the separate groups were not independent, it would be $Var(X + Y + Z) = Var(X) + Var(Y) + Var(Z) + 2 \cdot Cov(X, Y) + 2 \cdot Cov(X, Z) + 2 \cdot Cov(Y, Z)$. However, we are assuming that the variance of each of the groups is independent, so $Var(X + Y + Z) = Var(X) + Var(Y) + Var(Z)$. Recall $\psi = \mu_1 + \frac{1}{2}(\mu_2 + \mu_3)$. The variance would be

$$Var = Var(\bar{Y}_1) + Var(-\frac{1}{2}\bar{Y}_2) + Var(-\frac{1}{2}\bar{Y}_3)$$

$$= \sigma^2 \left(\frac{1}{n_1} + \frac{(-1/2)^2}{n_2} + \frac{(-1/2)^2}{n_3} \right)$$

$$E(S_p^2) = \sigma^2$$

$$sd = S_p \sqrt{\left(\frac{1}{n_1} + \frac{(-1/2)^2}{n_2} + \frac{(-1/2)^2}{n_3} \right)}$$

$$T = \frac{\bar{Y}_1 - \frac{1}{2}(\bar{Y}_2 + \bar{Y}_3)}{S_p \sqrt{\left(\frac{1}{n_1} + \frac{(-1/2)^2}{n_2} + \frac{(-1/2)^2}{n_3} \right)}}$$

$$S_p^2 = SSE/df_E \text{ (denom of } F \text{ test)}$$

Multiple Comparisons

Since contrast tests involves multiple comparisons, we have an inflated risk of type I errors. This can be corrected using the Bonferroni correction, where $\alpha^* = \alpha/(\text{number of tests})$. For N simultaneous tests that are independent, the overall type I error (ie. the probability of at least one type I error among N tests) is $1 - (1 - \alpha)^N$. For N simultaneous tests that are perfectly positively dependent (if one rejects, then the other combinations are

likely to reject as well, even if the groups are independent), the overall type I error is just α .

For N simultaneous tests that have unknown dependence, the maximum overall type I error is $\min(N\alpha, 1)$ (since it is a probability, it can't be > 1), where

$$\alpha \leq 1 - (1 - \alpha)^N \leq N\alpha$$

Tukey Honest Significant Difference (HSD)

This is the correct adjustment under ANOVA. Basic idea: consider the *largest difference* between any two sample means for for I groups.

Intuition: There are $\bar{Y}_1, \bar{Y}_2, \dots, \bar{Y}_I$ sample means. To do a pairwise comparison would generate $\binom{I}{2}$ 2 sample t -tests. We are going to reject one of these with 5% comparison rates. So it reduces to only have to check the ratio of max and min sample means.

It is appropriate when interested in differences between all pairs of group means. We assume normality, equal variances, equal sample sizes(\bar{n}), and equal means.

$$Q = \frac{\bar{Y}_{max} - \bar{Y}_{min}}{S_p \sqrt{\bar{n}}} \sim q(1 - \alpha, n - I), \text{ where } n = \bar{n}I$$

The Tukey distribution q is called the **Studentized Range Distribution**. In R, use `qtukey()` to obtain quantiles.

The margin of for the Confidence Interval (This is also used for the critical value as part of the t -test):

$$\left(\frac{q(1 - \alpha, I, n - I)}{\sqrt{2}} \right) SE, \text{ where } SE = S_p \sqrt{\frac{1}{\bar{n}} + \frac{1}{\bar{n}}}$$

Two-way and multi-way ANOVA

One way ANOVA - only one factor affecting the groups.

Two way ANOVA - multiple grouping variables to compare groups.

This can be reduced to boxplot

Interactive effect: when one variable changes the other variable. E.g. by going from sophomore to junior, the number of texts would change. degrees of freedom = $(I_1 - 1) * (I_2 - 1) = (4 - 1) * (2 - 1) = 3$

Kruskal-Wallis Test

This is a form of ANOVA on Ranks rather than t -test. The KW test is just an extension of the Wilcoxon Rank Sum test to 3 or more groups. The same procedure as other rank tests:

1. Rank all the combined data ignoring groups from 1 to N (treating them all like one sample). For any ties, average those ranks.
2. Calculate an F like χ^2 test statistic:

$$K = (N - 1) \frac{\sum_{i=1}^I n_i (\bar{R}_i - \bar{R})^2}{\sum_{i=1}^I \sum_{j=1}^{n_i} (R_{ij} - \bar{R})^2} \sim \chi_{I-1}^2$$

The hypothesis is like the Rank-Sum test: the quantiles in the group are all the same.

It is χ^2 because the numerator and denominator are tightly related. Only the numerator is a random variable.